

# **ECE 327/627**

# **Digital Hardware Systems**

## **Lecture 9: Latency Insensitive Design**

Andrew Boutros

[andrew.boutros@uwaterloo.ca](mailto:andrew.boutros@uwaterloo.ca)

# Some Logistics

- Lab 2 deadline is **tomorrow @ 11:59 pm**
  - Worth **8%** of the total ECE 327 grade
  - Worth **4%** of the total ECE 327 grade
- **[ECE 327]** Lab 2 assessment quiz in next week's lab sessions
  - Worth **40%** of the lab grade

# In the Previous Lecture ...

- Retiming
  - Algorithmic approach for improving frequency by moving around pipeline registers using simple rules
  - Preserves circuit behavior → does not add new registers
- Pipelining circuits with feedback loops
  - Possible but requires interleaving the computation of multiple streams of data to achieve correct functionality while having full throughput
- Throughput balancing in streaming architectures
  - Pipe and water analogy
  - Throughput of entire structure = the minimum throughput of all stages
  - Throughput balancing by module replication
  - Throughput balancing by multi-pumping

# Chip-wide Interconnect Roadblock

- In the good old days, the delay of a digital circuit could be considered the sum of delays of cascaded logic gates

# Chip-wide Interconnect Roadblock

- In the good old days, the delay of a digital circuit could be considered the sum of delays of cascaded logic gates
- This is no longer the case ...
  - In modern process generations, wire delay is becoming dominant
  - Wires get thinner and pitch (i.e., spacing) between wires get smaller
    - Resistance & Capacitance increase for a constant wire length

# Chip-wide Interconnect Roadblock

- In the good old days, the delay of a digital circuit could be considered the sum of delays of cascaded logic gates
- This is no longer the case ...
  - In modern process generations, wire delay is becoming dominant
  - Wires get thinner and pitch (i.e., spacing) between wires get smaller
    - Resistance & Capacitance increase for a constant wire length
- Some interesting numbers to put things in perspective
  - At 60nm process technology, a signal can reach only 5% of the chip's length in a clock cycle [1]
  - A 1mm interconnect at 35nm process technology is  $\sim 100x$  longer than a transistor switching delay [2]
  - On a 14nm FPGA, chip-spanning connection delay is  $\sim 10\text{ns}$  [3]

[1] D. Matzke, "Will Physical Scalability Sabotage Performance Gains?", IEEE Computer, 1997

[2] J. Davis, et al., "Interconnect Limits on Gigascale Integration in the 21<sup>st</sup> Century", Proceedings of the IEEE, 2001

[3] M. Abbas, V. Betz, "Latency Insensitive Design Styles for FPGAs", FPL, 2018

# Chip-wide Interconnect Roadblock

- In the good old days, the delay of a digital circuit could be considered the sum of delays of cascaded logic gates
- This is no longer the case ...
  - In modern process generations, wire delay is becoming dominant
  - Wires get thinner and pitch (i.e., spacing) between wires get smaller
    - Resistance & Capacitance increase for a constant wire length
- Some interesting numbers to put things in perspective
  - At 60nm process technology, a signal can reach only 5% of the chip's length in a clock cycle [1]
  - A 1mm interconnect at 35nm process technology is  $\sim 100x$  longer than a transistor switching delay [2]
  - On a 14nm FPGA, chip-spanning connection delay is  $\sim 10\text{ns}$  [3]
- Most problematic for “global” inter-module connections

[1] D. Matzke, “Will Physical Scalability Sabotage Performance Gains?”, IEEE Computer, 1997

[2] J. Davis, et al., “Interconnect Limits on Gigascale Integration in the 21<sup>st</sup> Century”, Proceedings of the IEEE, 2001


[3] M. Abbas, V. Betz, “Latency Insensitive Design Styles for FPGAs”, FPL, 2018

## Luckily, We Know a Trick ...

- Pipelining “global” wires to improve operating frequency is now a necessity

# Luckily, We Know a Trick ...


- Pipelining “global” wires to improve operating frequency is now a necessity
- We already know how to pipeline wires – it’s simple!



```
module wire_pipeline# (
    parameter DATAW = 8,
    parameter NUM_STAGES = 3
)(  
    input  clk,  
    input  rst,  
    input  [DATAW-1:0] data_in,  
    output [DATAW-1:0] data_out  
>);  
  
logic [DATAW-1:0] r_data [0:NUM_STAGES-1];  
integer i;  
  
always_ff @ (posedge clk) begin  
    if (rst) begin  
        for (i=0; i<NUM_STAGES; i=i+1) begin  
            r_data[i] <= 0;  
        end  
    end  
    else begin  
        r_data[0] <= data_in;  
        for (i=1; i<NUM_STAGES; i=i+1) begin  
            r_data[i] <= r_data[i-1];  
        end  
    end  
end  
assign data_out = r_data[NUM_STAGES-1];  
endmodule
```

# Luckily, We Know a Trick ...

- Pipelining “global” wires to improve operating frequency is now a necessity
- We already know how to pipeline wires – it’s simple!




- As many other things in life, turns out it is not as simple as it seems!

```
module wire_pipeline# (
    parameter DATAW = 8,
    parameter NUM_STAGES = 3
)(  
    input  clk,  
    input  rst,  
    input  [DATAW-1:0] data_in,  
    output [DATAW-1:0] data_out  
>;  
  
logic [DATAW-1:0] r_data [0:NUM_STAGES-1];  
integer i;  
  
always_ff @ (posedge clk) begin  
    if (rst) begin  
        for (i=0; i<NUM_STAGES; i=i+1) begin  
            r_data[i] <= 0;  
        end  
    end  
    else begin  
        r_data[0] <= data_in;  
        for (i=1; i<NUM_STAGES; i=i+1) begin  
            r_data[i] <= r_data[i-1];  
        end  
    end  
end  
assign data_out = r_data[NUM_STAGES-1];  
endmodule
```


# Problem with Pipelining Wires

Let's think about the simple case of communication between two modules on the chip ...



# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & exactly 2 cycles later use the data to perform some functionality

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & exactly 2 cycles later use the data to perform some functionality

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & exactly 2 cycles later use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & exactly 2 cycles later use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires

Let's think about the simple case of communication between two modules on the chip ...




- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & **exactly 2 cycles later** use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires

Let's think about the simple case of communication between two modules on the chip ...


*But now M1 should use data 4 cycles after issuing an address!*



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1 is designed to issue an address & **exactly 2 cycles later** use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- **M1.1** is designed to issue an address & **exactly 4 cycles later** use the data to perform some functionality

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- **M1.1** is designed to issue an address & **exactly 4 cycles later** use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- **M1.2** is designed to issue an address & **exactly 6 cycles later** use the data to perform some functionality
- Run through the implementation CAD tools → not meeting timing because of the wire delays between the two modules

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- M1.3 is designed to issue an address & **exactly 8 cycles later** use the data to perform some functionality

# Problem with Pipelining Wires


Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- **M1.3** is designed to issue an address & **exactly 8 cycles later** use the data to perform some functionality
- Now it's time to move to different process generation or a different FPGA device!!

# Problem with Pipelining Wires

Let's think about the simple case of communication between two modules on the chip ...



- Assume the memory block takes 2 cycles to output the data stored in a specific address
- **M1.3** is designed to issue an address & **exactly 8 cycles later** use the data to perform some functionality
- Now it's time to move to different process generation or a different FPGA device!!

# Latency-Insensitive Design (LID)

- Decouple compute functionality & communication




# Latency-Insensitive Design (LID)

- Decouple compute functionality & communication
- Communication between modules happen over “channels”
  - Channels can have arbitrary latencies (hence the name “latency-insensitive” channel)
  - Modules perform computation whenever **valid** inputs arrive, and they are **ready** to accept it
  - Regardless of when that happens, it does not affect functionality



# Latency-Insensitive Design (LID)


- Decouple compute functionality & communication
- Communication between modules happen over “channels”
  - Channels can have arbitrary latencies (hence the name “latency-insensitive” channel)
  - Modules perform computation whenever **valid** inputs arrive, and they are **ready** to accept it
  - Regardless of when that happens, it does not affect functionality



*A communication channel could be **simple wires***

# Latency-Insensitive Design (LID)


- Decouple compute functionality & communication
- Communication between modules happen over “channels”
  - Channels can have arbitrary latencies (hence the name “latency-insensitive” channel)
  - Modules perform computation whenever **valid** inputs arrive, and they are **ready** to accept it
  - Regardless of when that happens, it does not affect functionality



*A communication channel could be **pipelined wires***

# Latency-Insensitive Design (LID)


- Decouple compute functionality & communication
- Communication between modules happen over “channels”
  - Channels can have arbitrary latencies (hence the name “latency-insensitive” channel)
  - Modules perform computation whenever **valid** inputs arrive, and they are **ready** to accept it
  - Regardless of when that happens, it does not affect functionality



*A communication channel could be **shared bus***


# Latency-Insensitive Design (LID)

- Decouple compute functionality & communication
- Communication between modules happen over “channels”
  - Channels can have arbitrary latencies (hence the name “latency-insensitive” channel)
  - Modules perform computation whenever **valid** inputs arrive, and they are **ready** to accept it
  - Regardless of when that happens, it does not affect functionality




*A communication channel could be **network-on-chip (NoC)***

# Ready-Valid Handshake



Since there is no strict requirement on when data arrives ...


# Ready-Valid Handshake



Since there is no strict requirement on when data arrives ...

- Sender must be able to indicate that it is safe for the receiver to read and use the data → **Valid** bit


# Ready-Valid Handshake



Since there is no strict requirement on when data arrives ...

- Sender must be able to indicate that it is safe for the receiver to read and use the data → **Valid** bit *Similar to what we discussed when handling pipeline bubbles!*


# Ready-Valid Handshake



Since there is no strict requirement on when data arrives ...

- Sender must be able to indicate that it is safe for the receiver to read and use the data → **Valid** bit
- If receiver cannot accept data, sender must be notified to hold the message → **Ready** bit


# Ready-Valid Handshake



Since there is no strict requirement on when data arrives ...


- Sender must be able to indicate that it is safe for the receiver to read and use the data → **Valid** bit
- If receiver cannot accept data, sender must be notified to hold the message → **Ready** bit
- States of the handshake:
  - If valid == 1 & ready == 1 → transmission happens
  - If valid == 1 & ready == 0 → sender is holding the data
  - If valid == 0 & ready == X → there is nothing to send (channel is idle)

# Ready-Valid Handshake Expected Behavior




- At  $t=1$ , Sender produced data 'A' and tagged it as valid

# Ready-Valid Handshake Expected Behavior




- At t=1, Sender produced data 'A' and tagged it as valid
- At t=2, transmission of 'A' happens (valid && ready)

# Ready-Valid Handshake Expected Behavior




- At t=1, Sender produced data 'A' and tagged it as valid
- At t=2, transmission of 'A' happens (valid && ready)
  - Sender has no more data to send → valid is de-asserted
  - Receiver not ready to accept more data → ready is de-asserted

# Ready-Valid Handshake Expected Behavior




- At t=4, although receiver is not ready, sender produced data ‘B’ and tagged it as valid

# Ready-Valid Handshake Expected Behavior



- At  $t=4$ , although receiver is not ready, sender produced data ‘B’ and tagged it as valid
- At  $t=5$ , receiver became ready again!

# Ready-Valid Handshake Expected Behavior



- At t=4, although receiver is not ready, sender produced data ‘B’ and tagged it as valid – has to be held until receiver is ready
- At t=5, receiver became ready again!
- At t=6, transmission of ‘B’ happens (valid && ready)
  - Sender has more data ‘C’ but receiver is not ready to accept it

# If Receiver not Ready, Stall the Pipeline!



- A pipelined module (or a pipelined communication channel) has multiple pieces of data “in-flight” at different stages of the pipeline

# If Receiver not Ready, Stall the Pipeline!


- A pipelined module (or a pipelined communication channel) has multiple pieces of data “in-flight” at different stages of the pipeline
- If receiver is not ready, we need to freeze the pipeline in-place so data is not lost
  - This is known as “stalling” the pipeline

# If Receiver not Ready, Stall the Pipeline!


- A pipelined module (or a pipelined communication channel) has multiple pieces of data “in-flight” at different stages of the pipeline
- If receiver is not ready, we need to freeze the pipeline in-place so data is not lost
  - This is known as “stalling” the pipeline
- Registers typically have a clock enable (ce) signal
  - If  $ce = 1$ , input data is registered at clock edge
  - If  $ce = 0$ , register keeps its value (i.e., nothing happens at clock edge)




# Example: Latency Insensitive Polynomial Circuit




# Example: Latency Insensitive Polynomial Circuit




# Example: Latency Insensitive Polynomial Circuit




# Example: Latency Insensitive Polynomial Circuit




# Example: Latency Insensitive Polynomial Circuit



# Example: Latency Insensitive Polynomial Circuit



# Example: Latency Insensitive Polynomial Circuit



# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y

);

// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
```

```
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;

    end else begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;

        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;

        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;

        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;

        // Stage 4
        y <= r_y3 + r_c3;

    end
end

endmodule
```

# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y,
    input i_valid,
    input i_ready,
    output logic o_valid,
    output o_ready
);

// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
```

```
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;

    end else begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;

        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;

        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;

        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;

        // Stage 4
        y <= r_y3 + r_c3;

    end
end

endmodule
```

# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y,
    input i_valid,
    input i_ready,
    output logic o_valid,
    output o_ready
);

// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;
// Valid signal registers
logic r_valid, r_valid1, r_valid2, r_valid3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
```

```
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;

    end else begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;

        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;

        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;

        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;

        // Stage 4
        y <= r_y3 + r_c3;
    end
end

endmodule
```

# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y,
    input i_valid,
    input i_ready,
    output logic o_valid,
    output o_ready
);

// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;
// Valid signal registers
logic r_valid, r_valid1, r_valid2, r_valid3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
```

```
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;
        r_valid <= 0; r_valid1 <= 0;
        r_valid2 <= 0; r_valid3 <= 0;
    end else begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;
        r_valid <= i_valid;
        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;
        r_valid1 <= r_valid;
        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;
        r_valid2 <= r_valid1;
        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;
        r_valid3 <= r_valid2;
        // Stage 4
        y <= r_y3 + r_c3;
        o_valid <= r_valid3;
    end
end

endmodule
```

# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y,
    input i_valid,
    input i_ready,
    output logic o_valid,
    output o_ready
);

// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;
// Valid signal registers
logic r_valid, r_valid1, r_valid2, r_valid3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;
        r_valid <= 0; r_valid1 <= 0;
        r_valid2 <= 0; r_valid3 <= 0;
    end else begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;
        r_valid <= i_valid;
        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;
        r_valid1 <= r_valid;
        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;
        r_valid2 <= r_valid1;
        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;
        r_valid3 <= r_valid2;
        // Stage 4
        y <= r_y3 + r_c3;
        o_valid <= r_valid3;
    end
end

assign o_ready = i_ready;

endmodule
```


# Example: Latency Insensitive Polynomial Circuit

```
module poly_li (
    input clk,
    input rst,
    input [7:0] a,
    input [7:0] b,
    input [7:0] c,
    input [7:0] x,
    output logic [23:0] y,
    input i_valid,
    input i_ready,
    output logic o_valid,
    output o_ready
);


// Input registers
logic [7:0] r_a, r_b, r_c, r_x;
// Stage 1 registers
logic [15:0] r_y1;
logic [7:0] r_x1, r_b1, r_c1;
// Stage 2 registers
logic [15:0] r_y2;
logic [7:0] r_x2, r_c2;
// Stage 3 registers
logic [23:0] r_y3;
logic [7:0] r_c3;
// Valid signal registers
logic r_valid, r_valid1, r_valid2, r_valid3;

always_ff @ (posedge clk) begin
    if (rst) begin
        r_a <= 0; r_b <= 0;
        r_c <= 0; r_x <= 0;
        r_y1 <= 0; r_y2 <= 0; r_y3 <= 0;
        r_x1 <= 0; r_b1 <= 0; r_c1 <= 0;
        r_x2 <= 0; r_c2 <= 0;
        r_c3 <= 0; y <= 0;
        r_valid <= 0; r_valid1 <= 0;
        r_valid2 <= 0; r_valid3 <= 0;
    end else if (i_ready) begin
        // Input registers
        r_a <= a; r_b <= b;
        r_c <= c; r_x <= x;
        r_valid <= i_valid;
        // Stage 1
        r_y1 <= (r_a * r_x);
        r_x1 <= r_x; r_b1 <= r_b; r_c1 <= r_c;
        r_valid1 <= r_valid;
        // Stage 2
        r_y2 <= r_y1 + r_b1;
        r_x2 <= r_x1; r_c2 <= r_c1;
        r_valid2 <= r_valid1;
        // Stage 3
        r_y3 <= r_y2 * r_x2;
        r_c3 <= r_c2;
        r_valid3 <= r_valid2;
        // Stage 4
        y <= r_y3 + r_c3;
        o_valid <= r_valid3;
    end
    assign o_ready = i_ready;
endmodule
```

# Example: Latency Insensitive Polynomial Circuit




# Example: Latency Insensitive Polynomial Circuit



All pipeline  
registers are  
frozen!


# “Global Stalling” is Simple, but ...



# “Global Stalling” is Simple, but ...

There are two major issues with this approach ...

1. Ready signal has large fanout – it must connect to all clock enable inputs of all registers
2. Ready signal must be purely combinational




# “Global Stalling” is Simple, but ...


There are two major issues with this approach ...

1. Ready signal has large fanout – it must connect to all clock enable inputs of all registers
2. Ready signal must be purely combinational

*Why can't  
we just  
pipeline the  
ready signal  
going  
backwards?*




# Well ... That didn't solve our problem!



- We can pipeline the data & valid signals as much as we want without needing to change the internal logic of the receiver
- However, we cannot pipeline the ready signal to avoid losing data in case of a backpressure

***Simple pipeline registers are not enough!!***


# Relay Station



# Relay Station

Consists of ...


- Main and auxiliary registers for data and valid signals



# Relay Station

Consists of ...


- Main and auxiliary registers for data and valid signals
- MUXEs to allow writing to main registers from inputs or aux registers



# Relay Station


Consists of ...

- Main and auxiliary registers for data and valid signals
- MUXes to allow writing to main registers from inputs or aux registers
- Control FSM to issue register clock enables and MUX select lines
  - Remember – the FSM has registered outputs




# Relay Station

- If receiver is ready, the main registers are used to propagate the data and valid signals




# Relay Station

- If receiver is not ready, the main registers are stalled
- Since  $i_{ready}$  takes 1 cycle to propagate to  $o_{ready}$ , sender has already sent a new data/valid, hold them in aux registers




# Relay Station

- When receiver is ready again, first read out data/valid from main registers and push data/valid from aux registers to main




# Relay Station


- When receiver is ready again, first read out data/valid from main registers and push data/valid from aux registers to main
- Then back to normal operation ...




# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Example Operation of Relay Stations



# Example Operation of Relay Stations




# Example Operation of Relay Stations




# Relay Station Control FSM

- FSM has 4 states ...
  - **PROCESS**: normal operation
  - **WR\_AUX**: write to aux reg when  $i_{ready}$  goes from  $1 \rightarrow 0$
  - **RD\_AUX**: read from aux reg when  $i_{ready}$  goes from  $0 \rightarrow 1$
  - **STALL**: relay station is stalled



# Relay Station Control FSM


- FSM has 4 states ...
  - **PROCESS**: normal operation
  - **WR\_AUX**: write to aux reg when  $i_{ready}$  goes from  $1 \rightarrow 0$
  - **RD\_AUX**: read from aux reg when  $i_{ready}$  goes from  $0 \rightarrow 1$
  - **STALL**: relay station is stalled
- Transition between states based on value of  $i_{ready}$



# Relay Station Control FSM


Set MUX select lines and register clock enables:

- If writing to aux regs or stalling, freeze main regs ( $\text{main\_ce} = 0$ ), otherwise clock enable them ( $\text{main\_ce} = 1$ )
- If writing to aux regs, clock enable them ( $\text{aux\_ce} = 1$ ), otherwise freeze them ( $\text{aux\_ce} = 0$ )
- MUX select chooses aux register as a source if stalling or coming out of a stall ( $\text{mux\_sel} = 1$ ), otherwise chooses input data as a source ( $\text{mux\_sel} = 0$ )



# Relay Station SystemVerilog Code


```
module relay_station # (
    parameter DATAW = 32
) (
    input clk,
    input rst,
    input [DATAW-1:0] i_data,
    input i_valid,
    input i_ready,
    output logic [DATAW-1:0] o_data,
    output logic o_valid,
    output logic o_ready
);
```



# Relay Station SystemVerilog Code

```
module relay_station # (
    parameter DATAW = 32
) (
    input clk,
    input rst,
    input [DATAW-1:0] i_data,
    input i_valid,
    input i_ready,
    output logic [DATAW-1:0] o_data,
    output logic o_valid,
    output logic o_ready
);

logic [DATAW-1:0] data_aux;
logic valid_aux;
enum {PROCESS, WR_AUX, RD_AUX, STALL} state, next_state;
logic main_ce_w, aux_ce_w, mux_sel_w;
```



# Relay Station SystemVerilog Code


```

module relay_station # (
    parameter DATAW = 32
)
(
    input clk,
    input rst,
    input [DATAW-1:0] i_data,
    input i_valid,
    input i_ready,
    output logic [DATAW-1:0] o_data,
    output logic o_valid,
    output logic o_ready
);

logic [DATAW-1:0] data_aux;
logic valid_aux;
enum {PROCESS, WR_AUX, RD_AUX, STALL} state, next_state;
logic main_ce_w, aux_ce_w, mux_sel_w;

always_ff @ (posedge clk) begin
    if (rst) begin
        o_data <= 'd0; data_aux <= 'd0;
        o_valid <= 1'b0; valid_aux <= 1'b0;
    end else begin
        if (main_ce_w) begin
            o_data <= (mux_sel_w)? data_aux : i_data;
            o_valid <= (mux_sel_w)? valid_aux : i_valid;
        end
        if (aux_ce_w) begin
            data_aux <= i_data;
            valid_aux <= i_valid;
        end
    end
end
end

```



# Relay Station SystemVerilog Code

```

module relay_station # (
    parameter DATAW = 32
)
(
    input clk,
    input rst,
    input [DATAW-1:0] i_data,
    input i_valid,
    input i_ready,
    output logic [DATAW-1:0] o_data,
    output logic o_valid,
    output logic o_ready
);

logic [DATAW-1:0] data_aux;
logic valid_aux;
enum {PROCESS, WR_AUX, RD_AUX, STALL} state, next_state;
logic main_ce_w, aux_ce_w, mux_sel_w;


always_ff @ (posedge clk) begin
    if (rst) begin
        o_data <= 'd0; data_aux <= 'd0;
        o_valid <= 1'b0; valid_aux <= 1'b0;
    end else begin
        if (main_ce_w) begin
            o_data <= (mux_sel_w)? data_aux : i_data;
            o_valid <= (mux_sel_w)? valid_aux : i_valid;
        end
        if (aux_ce_w) begin
            data_aux <= i_data;
            valid_aux <= i_valid;
        end
    end
end
end

```

```

always_ff @ (posedge clk) begin
    if (rst) begin
        state <= PROCESS;
        o_ready <= 1'b0;
    end else begin
        state <= next_state;
        o_ready <= i_ready;
    end
end

```



# Relay Station SystemVerilog Code

```

module relay_station # (
    parameter DATAW = 32
)
(
    input clk,
    input rst,
    input [DATAW-1:0] i_data,
    input i_valid,
    input i_ready,
    output logic [DATAW-1:0] o_data,
    output logic o_valid,
    output logic o_ready
);

logic [DATAW-1:0] data_aux;
logic valid_aux;
enum {PROCESS, WR_AUX, RD_AUX, STALL} state, next_state;
logic main_ce_w, aux_ce_w, mux_sel_w;

always_ff @ (posedge clk) begin
    if (rst) begin
        o_data <= 'd0; data_aux <= 'd0;
        o_valid <= 1'b0; valid_aux <= 1'b0;
    end else begin
        if (main_ce_w) begin
            o_data <= (mux_sel_w)? data_aux : i_data;
            o_valid <= (mux_sel_w)? valid_aux : i_valid;
        end
        if (aux_ce_w) begin
            data_aux <= i_data;
            valid_aux <= i_valid;
        end
    end
end
end


```

```

always_ff @ (posedge clk) begin
    if (rst) begin
        state <= PROCESS;
        o_ready <= 1'b0;
    end else begin
        state <= next_state;
        o_ready <= i_ready;
    end
end

always_comb begin: state_decoder
    case (state)
        PROCESS: next_state = (!i_ready)? WR_AUX : PROCESS;
        WR_AUX: next_state = (!i_ready)? STALL : RD_AUX;
        RD_AUX: next_state = (!i_ready)? WR_AUX : PROCESS;
        STALL: next_state = (!i_ready)? STALL : RD_AUX;
        default: next_state = PROCESS;
    endcase
end

```




# Relay Station SystemVerilog Code

```


always_comb begin: output_decoder
    case (state)
        PROCESS: begin
            mux_sel_w = 1'b0;
            main_ce_w = (!i_ready)? 1'b0 : 1'b1;
            aux_ce_w = (!i_ready)? 1'b1 : 1'b0;
        end
        WR_AUX: begin
            mux_sel_w = 1'b1;
            main_ce_w = (!i_ready)? 1'b0 : 1'b1;
            aux_ce_w = 1'b0;
        end
        RD_AUX: begin
            mux_sel_w = 1'b0;
            main_ce_w = (!i_ready)? 1'b0 : 1'b1;
            aux_ce_w = (!i_ready)? 1'b1 : 1'b0;
        end
        STALL: begin
            mux_sel_w = 1'b1;
            main_ce_w = (!i_ready)? 1'b0: 1'b1;
            aux_ce_w = 1'b0;
        end
        default: begin
            mux_sel_w = 1'b0;
            main_ce_w = (!i_ready)? 1'b0 : 1'b1;
            aux_ce_w = (!i_ready)? 1'b1 : 1'b0;
        end
    end
endcase
end

endmodule

```



# Relay Stations Save the Day!



- We can add an arbitrary number of relay stations between sender and receiver
  - No combinational paths → easier to close timing!
  - No need to change logic in sender/receiver as we change number of relay stations on the communication channel
- Relay stations are sometimes referred to as “skid buffers” or “latency-insensitive register slice”

# Advanced eXtensible Interface (AXI) Protocol


- AXI is a latency-insensitive interface protocol
  - Instead of the simple data/valid/ready interface we have been using so far, AXI defines different types of interfaces for different use cases
  - Introduced by ARM in 2003 as part of the Advanced Microcontroller Bus Architecture (AMBA) standard
  - The industry standard for latency-insensitive communication between different modules in a digital system
  - Full specification can be found at this [link](#)
- Designing an intellectual property (IP) block that speaks AXI enables others to easily integrate it in their systems
- For FPGA development, AMD Xilinx adopted AXI while Altera developed a similar protocol called Avalon – now gradually shifting to AXI for wider compatibility

# Types of AXI Interfaces


There are 3 types of AXI interfaces used for different purposes ...

1. AXI Memory Mapped (typically referred to as just “AXI”)
  - Used to read/write data to/from a specific address
  - Has 5 channels: read address (AR), read data (R), write address (AW), write data (W), and write response (B)
2. AXI Lite
  - Simplified version of AXI Memory Mapped
  - Used for simple, low-throughput, memory-mapped communication such as reading or writing from Control and Status Registers (CSRs)
3. AXI Stream
  - Used for data flowing/streaming between two end points
  - No notion of addresses

# Types of AXI Interfaces




# Types of AXI Interfaces




The AXI interface initiating a transaction is called the **Manager** and the AXI interface receiving or responding to a transaction is called the **Subordinate**

# AXI Stream Interface Signals




# AXI Memory Mapped



Source: [amd.com](http://amd.com)

# AXI Memory Mapped



## Next Lecture ...

FIFO-based design ...